Limitations

- Currently doesn't track why invalidation occurred leading to unnecessary rework

- Hard to delete (garbage collect) old values

- Running function generators probably corrupts inference due to mis-ordering when reloading

- Incremental restoration algorithm perhaps still not entirely correct

- Doesn't require on-stack replacement

Questions?