

~~Statically~~ Compiling Julia

Jameson Nash

Julia is an Ahead-of-Time (AoT) statically-compiled language*!

*(*excludes staged functions)*

But doesn't it use a JIT?

The homepage of Julia describes the LLVM JIT compiler that helps to make Julia fast.

What it doesn't mention, is that Julia's language design permits better Ahead-of-Time analysis than a traditional JIT must deal with.

- Elegant and extensible [conversions and promotions](#) for
- Efficient support for [Unicode](#), including but not limited
- [MIT licensed](#): free and open source

High-Performance JIT Compiler

Julia's LLVM-based just-in-time (JIT) compiler combined v often match the performance of C. To get a sense of relative that can or could be used for numerical and scientific comp benchmarks in a variety of languages: [C](#), [Fortran](#), [Julia](#), [Pyt](#) and [Mathematica](#). We encourage you to skim the code to ge

(<http://julialang.org>)

What does AoT mean?

Full pre-analysis (type inference) valid

Methods can't be changed on-stack

No runtime checks on types required by the JIT

Infamous issue #265: previous statements aren't reliably enforced by the runtime / compiler :(

Adding a method requires updating the AoT compiled code:

- “toplevel” code vs. function code
- natural separation: defining code vs. running data

Julia: REPL or Compiler?

```
~/julia$ ./julia --help
julia [switches] -- [programfile] [args...]
-v, --version          Display version information
-h, --help             Print this message

-J, --sysimage <file>  Start up with the given system image file
--precompiled={yes|no} Use precompiled code from system image if available
--compilecache={yes|no} Enable/disable incremental precompilation of modules
-H, --home <dir>       Set location of julia executable
--startup-file={yes|no} Load ~/.juliarc.jl
--handle-signal={yes|no} Enable or disable Julia's default signal handlers

-e, --eval <expr>       Evaluate <expr>
-E, --print <expr>      Evaluate and show <expr>
-L, --load <file>       Load <file> immediately on all processors

--compile={yes|no|all|min} Enable or disable JIT compiler, or request exhaustive compilation
-C, --cpu-target <target> Limit usage of cpu features up to <target>
-O, --optimize={0,1,2,3} Set the optimization level (default 2 if unspecified or 3 if specified as -O)
--inline={yes|no}        Control whether inlining is permitted (overrides functions declared as @inline)
--check-bounds={yes|no}  Emit bounds checks always or never (ignoring declarations)
--math-mode={ieee,fast}  Disallow or enable unsafe floating point optimizations (overrides @fastmath declaration)

--depwarn={yes|no|error} Enable or disable syntax and method deprecation warnings ("error" turns warnings into errors)

--output-o name          Generate an object file (including system image data)
--output-ji name         Generate a system image data file (.ji)
--output-bc name         Generate LLVM bitcode (.bc)
--output-incremental=no  Generate an incremental output file (rather than complete)

--code-coverage={none|user|all}, --code-coverage
                        Count executions of source lines (omitting setting is equivalent to "user")
--track-allocation={none|user|all}, --track-allocation
                        Count bytes allocated by each source line
```

Both?

-> Issue #15864: separate `julia-compile` from `julia`

Obligatory Word-cloud

Static

Compiler

Inlining

LLVM

HELP!

Parser

Speed

Source Code

Optimizations

Efficiency

Codegen

System
Image

Native code

Pre-
compilation

Macros

Runtime
Library

Modules

Embedding

Deployment

Mini Julia

> Is there some easy hackable mini-version of Julia available? #16798

Not easily. Julia isn't designed to be a “toy” language.

It is possible to simulate parts of Julia inside Julia:

<https://github.com/JuliaLang/julia/blob/master/examples/juliatypes.jl> : See Jeff's JuliaCon 2015 talk

But Julia is easily hackable! Let's go...

Main.Core

```
~/julia$ echo > nothing.jl
```

```
~/julia$ ./julia --output-ji minicore.ji \  
> nothing.jl
```

```
ERROR: could not open file boot.jl
```

```
~/julia$ cd base
```

```
~/julia/base$ ../julia --output-ji minicore.ji \  
> ../nothing.jl
```

```
~/julia/base$ ls -lh minicore.ji
```

```
-rw-r--r-- jameson 33K Jun 14 14:54 minicore.ji
```


REPL?

```
~/julia/base$ ../julia -J minicore.ji  
WARNING: Base._start not defined, falling  
back to economy mode repl.
```

```
julia> Core.print("Hello world\n")  
Hello world  
nothing
```

```
julia>
```

REPL?

```
~/julia/base$ ../julia -J minicore.ji  
WARNING: Base._start not defined, falling  
back to economy mode repl.
```

```
julia> 1 + 1  
error during run:  
UndefVarError(var=:+)
```

```
julia> Base  
error during run:  
UndefVarError(var=:Base)
```

Going beyond Core

```
~/julia/base$ ../julia --output-jl inference.jl coreimg.jl
```

```
~/julia/base$ ../julia --output-o sys.o --sysimage inference.jl sysimg.jl
```

```
~/julia/base$ cc -shared -o sys.so sys.o -ljulia -L ../usr/lib
```

```
~/julia/base$ ../julia --sysimage sys.so
```

```
      _
     _(_)_
    (_)|(_)(_)
   _ _ _|_|_ _ _
  | | | | | | | / _ ` |
  | | | _ | | | | (_ | |
 _/ | \ _ _ ' _ | _ | \ _ _ ' _ |
| _ _ /
      |
      | A fresh approach to technical computing
      | Documentation: http://docs.julialang.org
      | Type "?help" for help.
      |
      | Version 0.5.0-dev+4711 (2016-06-14 03:43 UTC)
      | jn/demo/09e75d4* (fork: 4 commits, 1 day)
      | x86_64-apple-darwin15.5.0
```

```
julia>
```

Deeper Inspection

But what did we just do?

- Julia as a compiler for writing Julia programs in Julia
- Static output with dynamic semantics

Overview: Translation stages

- `./julia (repl.c)`
- `eval/include (builtins.c)`
 - `parser (julia-parser.scm)`
 - `run macros (ast.c)`
 - `lowered form (julia-syntax.scm)`
 - `execute toplevel expressions (toplevel.c)`
 - `call functions (gf.c)`
 - `dispatch (typemap.c)`
 - `interpreter (interpreter.c)` OR
`compiler (codegen.cpp)`
+ `optimizations (inference.jl)`

For details, see Jeff's talk from JuliaCon 2014.

Overview: Standard Library

- `./julia (repl.c)`
 - getopt parsing
 - loads the standard library
- `libjulia.so` (or `.dll` or `.dylib`)
 - Garbage Collector / Allocator
 - Builtin Functions / Ininsics
 - Interpreter
 - Builtin Types
 - Type System
 - Dispatch System
 - Compiler
 - Platform abstraction / integration
 - System Image / serializer
- `sys.so` (and other `lib/julia/` files `*.so` and `*.ji`)
 - Core module
 - Inference module
 - Base module
- Packages
- Dependencies
 - LLVM, BLAS, libgit2, etc.

Minimal 'main(void)'

```
// This file (repl.c) is a part of Julia. License is MIT: http://julialang.org/license

// Standard headers
#include <string.h>
#include <stdint.h>
// Julia headers (for initialization and gc commands)
#include "uv.h"
#include "julia.h"

int main(int argc, char *argv[])
{
    // Initialize Julia
    uv_setup_args(argc, argv);
    libsupport_init();
    // Setup any runtime options here
    jl_options.fast_math = JL_OPTIONS_FAST_MATH_OFF; // --math-mode=ieee
    julia_init(JL_IMAGE_JULIA_HOME);
    // Run the REPL
    jl_function_t *start_client = (jl_function_t*)jl_get_global(jl_base_module, jl_symbol("_start"));
    (void)jl_apply(&start_client, 1);
    // Cleanup and graceful exit
    jl_atexit_hook(0);
    return 0;
}
```


More Standard Library

coreimg.jl

```
Main.Core.eval(Main.Core, :(baremodule Inference
using Core.Intrinsics
import Core: print, println, show, write, unsafe_write, STDOUT,
STDERR
ccall(:jl_set_istopmod, Void, (Bool,), false)

eval(x) = Core.eval(Inference,x)
eval(m,x) = Core.eval(m,x)

include = Core.include

## Load essential files and libraries
include("essentials.jl")
include("generator.jl")
include("reflection.jl")
include("options.jl")

# core operations & types
typealias Cint Int32
typealias Csize_t UInt
include("promotion.jl")
include("tuple.jl")
include("range.jl")
include("expr.jl")
include("error.jl")

# core numeric operations & types
include("bool.jl")
include("number.jl")
include("int.jl")
include("operators.jl")
include("pointer.jl")
const checked_add = +
const checked_sub = -
if !isdefined(Main, :Base)
    # conditional to allow redefining Core.Inference after base exists
```

sysimg.jl

```
baremodule Base
using Core.Intrinsics
ccall(:jl_set_istopmod, Void, (Bool,), true)
include = Core.include
include("coreio.jl")

eval(x) = Core.eval(Base,x)
eval(m,x) = Core.eval(m,x)

include("exports.jl")

## Load essential files and libraries
include("essentials.jl")
include("base.jl")
include("generator.jl")
include("reflection.jl")
include("options.jl")

# core operations & types
include("promotion.jl")
include("tuple.jl")
include("range.jl")
include("expr.jl")
include("error.jl")

# core numeric operations & types
include("bool.jl")
include("number.jl")
include("int.jl")
include("operators.jl")
include("pointer.jl")
include("refpointer.jl")
(::Type{T}){T}(arg) = convert(T, arg)::T
include("checked.jl")
importall .Checked
```

But not done yet...

using.jl

`--compile-incremental=yes`

using MyPackage

- Base.require(:MyPackage)
- isdefined(Main, :MyPackage) ?
 - Base._include_from_serialized or Base.include
 - Base.cachecompile(:MyPackage)
 - Recursion!
 - MyPackage.__init__()
- symbols are imported from MyPackage

Deployment

Future work in progress.

BuildExecutable.jl and contrib/build_sysimg.jl

```
int main(int argc, char *argv[ ]);
```

Output C code

Bundling incremental compile into application

Installer tarball

```
Project.generate( )
```



Profiling

```
~/julia/test$ ../julia runtests.jl core
```

```
    * core
```

```
  SUCCESS
```

| | | | |
|-----------------------|---|--------|-------------|
| REMAINDER | : | 8.50% | 5312822108 |
| GC | : | 13.15% | 8218142751 |
| LOWERING | : | 31.01% | 19385805839 |
| PARSING | : | 0.73% | 456448707 |
| INFERENCE | : | 11.16% | 6975897589 |
| CODEGEN | : | 1.28% | 799324355 |
| METHOD_LOOKUP_SLOW | : | 0.15% | 93144624 |
| METHOD_LOOKUP_FAST | : | 5.70% | 3566210459 |
| LLVM_OPT | : | 25.35% | 15844106508 |
| LLVM_MODULE_FINISH | : | 0.86% | 534538848 |
| METHOD_LOOKUP_COMPILE | : | 0.02% | 10894604 |
| TYPE_LOOKUP | : | 2.10% | 1315693342 |

Memory Usage

```
$ size -A sys.dylib
```

```
sys.dylib :
```

| section | size | addr |
|-----------------|----------|----------|
| __text | 4382405 | 3920 |
| __stubs | 366 | 4386326 |
| __stub_helper | 626 | 4386692 |
| __const | 26109544 | 4387328 |
| __cstring | 13108 | 30496880 |
| __unwind_info | 68376 | 30509988 |
| __eh_frame | 1251632 | 30578368 |
| __nl_symbol_ptr | 16 | 31830016 |
| __got | 64 | 31830032 |
| __la_symbol_ptr | 488 | 31830096 |
| __const | 406232 | 31830592 |
| __bss | 167488 | 32236824 |
| Total | 32400345 | |

```
~/julia/usr/lib$ size -A libjulia.dylib
```

```
libjulia.dylib :
```

| section | size | addr |
|-----------------|---------|---------|
| __text | 1098861 | 3424 |
| __stubs | 4350 | 1102286 |
| __stub_helper | 7166 | 1106636 |
| __cstring | 57318 | 1113808 |
| __const | 240682 | 1171136 |
| __unwind_info | 5416 | 1411820 |
| __eh_frame | 110560 | 1417240 |
| __nl_symbol_ptr | 16 | 1527808 |
| __got | 488 | 1527824 |
| __la_symbol_ptr | 5800 | 1528312 |
| __mod_init_func | 32 | 1534112 |
| __mod_term_func | 16 | 1534144 |
| __const | 327272 | 1534160 |
| __data | 219760 | 1861440 |
| __thread_vars | 24 | 2081200 |
| __thread_bss | 2368 | 2081224 |
| __bss | 806656 | 2083600 |
| __common | 1346169 | 2890256 |
| Total | 4232954 | |

POWER8 Support

[https://www-356.ibm.com/partnerworld/wps/ent/pdp/
web/MyProgramAccess](https://www-356.ibm.com/partnerworld/wps/ent/pdp/web/MyProgramAccess)

Future Work

Disable JIT

Better generic code

Incremental linking

Implicit main function support

Transparent shlib embedding

System image compression / “strip”

Extract much of Base into a standard package library

Ports: webassembly, microkernel, embedded

Fix #265 (but disable fix for release code / static compile?)

Further addressing package load times?

Native code caching archive database

Eliminate Core.Intrinsics (move to all BuiltinFunctions)

Questions?

<http://juliacomputing.com/blog/2016/02/09/static-julia.html>